

Youssef Ashraf

LinkedIn | GitHub | yousefashraf970@gmail.com | (+20) 01154202976

SUMMARY

Computer science professional passionate about deep learning and machine learning, with expertise in engineering AI-powered solutions. Experienced in building conversational AI systems, NLP pipelines, and deep learning models. Eager to contribute to innovative data-driven projects and continue deepening expertise in AI and data analysis.

EDUCATION

BNU (Benha National University) – Cairo, Egypt

B.Sc. in Computer Science, Major: Artificial Intelligence | Oct 2022 – July 2026 | GPA: 3.76

Graduated: July 2026 | Class Rank: 4th, AI Major

Relevant Coursework: Machine Learning · Deep Learning · Natural Language Processing · Neural Networks · Data Structures & Algorithms · Database Programming · Data Engineering & Visualization · Software Engineering · Statistics · Computer Networks

WORK EXPERIENCE

Cloud Services Trainee

2025

National Telecommunication Institute (NTI)

- Hands-on with 8 core AWS services: EC2, S3, IAM, VPC, Lambda, RDS, CloudWatch.
- Designed secure, scalable architectures aligned with Solutions Architect best practices.
- Reduced projected cloud costs by 20% through optimized IAM policies and networking.

Computer Vision Intern

2025

National Academy of IT for Persons with Disabilities (NAID) – New Capital · On-Site

- Contributed to a computer vision project for dyslexia detection; processed and analyzed 500+ text samples to improve reading assistance accuracy.
- Collaborated with a team of 6 researchers; contributed to 2 technical presentations on AI applications in accessibility, achieving 85% model accuracy in preliminary testing.

PROJECTS

BNUConnect – AI-Powered Academic Advisor Chatbot – Graduation Project

2026

Private Repo

- Engineered a production-grade AI academic advisor chatbot for Benha National University, deployed as a REST API (FastAPI + Uvicorn) integrated into the BNUConnect mobile app, serving students across 3 program tracks (AI & ML, Software Development, Data Science).
- Built a stateful LangGraph ReAct agent with a custom Judging Loop, orchestrating 12 LangChain tools across Neo4j knowledge graph (course catalog, prerequisites, program tracks), Supabase (student profiles, chat history), and a Pinecone RAG pipeline for university bylaw Q&A – achieving accurate multi-hop reasoning over the full academic domain.
- Designed a 5-step NLP preprocessing pipeline (reference resolution, entity extraction, fuzzy course name mapping, track normalization, query rewriting) that resolves abbreviations and pronouns before the agent runs, eliminating wasted tool calls and handling ambiguous queries via interactive disambiguation sessions.
- Implemented a RAG pipeline using HuggingFace multilingual-e5-large embeddings and Pinecone vector store (bnu-bylaws index) for bylaw/regulation Q&A, and built an interactive multi-turn course planner bridged from a CLI function to the chatbot via a thread I/O bridge with 90-second timeout handling.

Stack: LangGraph, LangChain, FastAPI, Groq, Neo4j Aura, Supabase (PostgreSQL), Pinecone, HuggingFace, Python, Docker.

BNU Admission Chatbot – Multilingual RAG Assistant

2026

Private Repo

- Built a multilingual RAG chatbot answering admission queries in Arabic (MSA & Egyptian), Franco-Arabic, and English, using hybrid dense (Pinecone) + sparse (BM25) retrieval with reciprocal rank fusion, plus an intent router that skips LLM calls for greetings, FAQs, and out-of-scope messages.
- Automated knowledge-base upkeep with a Google Drive sync job detecting new/modified/deleted source PDFs every 10 minutes, shipped as a FastAPI service with an RTL/LTR-aware frontend and Docker deployment.

Stack: FastAPI, Groq, Pinecone, BM25, Cohere, HuggingFace, Google Drive API, Docker.

Clash of Codes – Multilingual Historical RAG Chatbot – 1st Place, NLP Track

2026

Competition Project

- Built a RAG-based chatbot under competition conditions for Clash of Codes, an AI hackathon hosted by the Higher Technical Institute (HTI) on the NLP track, retrieving relevant historical-data passages before generating responses to keep answers factually grounded.
- Made the assistant trilingual – responding fluently in Arabic, English, and Franco-Arabic – and placed 1st among NLP-track teams at the competition.

Stack: RAG, NLP, Arabic NLP.

Auto Correction System

2025

Private Repo

- Designed and implemented a natural language processing system to automatically correct spelling and grammar errors using Transformer models and custom tokenization strategies.

Video Action Recognition on UCF50

2025

Public Repo

- Trained a spatio-temporal TensorFlow/Keras classifier to recognize 50 human action categories from the UCF50 dataset, sampling, resizing, and normalizing 20 frames per clip for model input.
- Deployed an interactive Gradio app that accepts uploaded videos, transcodes them with FFmpeg, and returns top-3 predicted actions with confidence scores.

Stack: TensorFlow, Keras, OpenCV, Gradio, FFmpeg, NumPy.

Twitter Sentiment Analysis with Apache Spark & Docker

2025

Public Repo

- Built a distributed sentiment analysis pipeline on a 3-node (1 master + 2 workers) Dockerized Apache Spark cluster, extracting 20,000 TF-IDF features and training a Logistic Regression classifier on ~75K tweets to label sentiment as Positive, Negative, Neutral, or Irrelevant (91.7% training / 80.2% validation accuracy).
- Served the model through a Flask web app with a REST endpoint for real-time predictions and per-class confidence scores.

Stack: Apache Spark, PySpark MLlib, Flask, Docker Compose, Logistic Regression, TF-IDF.

Face Recognition – Identity-Verified Account Access

2024

Public Repo

- Built a facial-recognition identity verification system using FaceNet (128-d embeddings) and MTCNN for face detection, matching users via L2 distance against 30+ stored identities.
- Wrapped the pipeline in a Gradio app simulating a secure bank-account flow (face login, account creation, deposit/withdraw) and deployed it live on HuggingFace Spaces.

Stack: TensorFlow, FaceNet, MTCNN, Gradio, NumPy.

Animal Segmentation Model with U-Net & TensorFlow

2024

Public Repo

- Built an animal segmentation model using the U-Net architecture with TensorFlow for pixel-wise classification. Involved preprocessing annotated datasets and evaluating performance with IoU and Dice coefficient metrics.

ACHIEVEMENTS

- Clash of Codes (HTI AI Hackathon) – 1st Place, NLP Track:** Placed 1st on the NLP track for a multilingual (Arabic, English, Franco-Arabic) RAG chatbot built over historical data under competition conditions.
- ECPC (Egyptian Collegiate Programming Contest) – Finalist:** Sole qualifying team from Benha National University; solved 4 problems as a 3-person team to advance from the qualification round to the ECPC finals (ACPC qualifier, under ICPC rules).

CERTIFICATIONS

- Cloud Architect – National Telecommunication Institute (NTI)
- Machine Learning Specialization – Andrew Ng (Coursera)
- Deep Learning Specialization – Andrew Ng (Coursera)
- Hugging Face Agents Course
- Clash of Codes – 1st Place, NLP Track (Higher Technical Institute)

TECHNICAL SKILLS

Programming Languages: Python, C++

AI / ML Frameworks: PyTorch, TensorFlow, Hugging Face Transformers, Scikit-learn

LLM / Agent Frameworks: LangChain, LangGraph, Smolagents, LlamaIndex, LangFuse

Databases & Vector Stores: Neo4j, Supabase (PostgreSQL), Pinecone

APIs & Tools: FastAPI, Groq, Gradio, Weights & Biases, GitHub

Deep Learning Techniques: CNNs, RNNs, LSTMs, Transformers, Transfer Learning, LLMs

Data Science Skills: Statistical Analysis, Data Visualization, Feature Engineering